



Torch - Scientific computing for Lua(JIT)

Octavian Ganea, Aurelien Lucchi

CLS Deep Learning Workshop

July 12th, 2016

Why use Torch ?

- ▶ Big community
 - ▶ Torch exists since 2000
 - ▶ FB, Twitter, IBM, Google, many universities and major labs
- ▶ Well documented
- ▶ Easy to deploy large-scale machine learning applications
 - ▶ Fast optimized backend (C, C++, Blas, CUDA)
 - ▶ Easy to run on GPU



Why use Torch ?

- ▶ Powerful
 - ▶ Easy to construct very complicated neural networks
 - ▶ lego-like base modules (network atomic ops) with automatic differentiation
 - ▶ Easy to use/modify/write any optimization method, layer, etc
 - ▶ Dynamic neural networks, unlike static ones in TF or Theano
 - ▶ best suited for research

Based on scripting language Lua

- ▶ Very easy to learn (Matlab like)
- ▶ Dense code - simple, readable, clean constructs
- ▶ Very easy to run on GPU
- ▶ Interpreted code, no time lost for compilation
- ▶ Easy portable code to other platform (e.g. iPhone)

- ▶ Interactive console



```
th>  | Torch7  
Scientific computing for Lua.  
Type ? for help  
https://github.com/torch  
http://torch.ch
```

Lua basics: tables

- ▶ Unique universal data structure
- ▶ Can be used as:
 - ▶ array / list
 - ▶ record/object
 - ▶ dictionary / hash table

```
th> a = {1, 3.25, 'ML', nil, true, false, {4,5,6}} -- (indexed) array
[0.0000s]
th> a
{
  1 : 1
  2 : 3.25
  3 : "ML"
  4 : true
  5 : false
  6 : {
    1 : 4
    2 : 5
    3 : 6
  }
}
[0.0001s]
th> a = {x=0, y=0, label="console"} -- record/object
[0.0000s]
th> a
{
  label : "console"
  y : 0
  x : 0
}
[0.0001s]
th> a.label
console
[0.0000s]
th> a['label']
console
[0.0000s]
th> a.x = 1.0
[0.0000s]
th>
[0.0000s]
th> a = {[ "+" ] = "add", [ "-" ] = "sub", [ "*" ] = "mul", [ "/" ] = "div"} -- dictionary / map
[0.0000s]
th> for key,value in pairs(a) do
..> print(key .. ' -> ' .. value)
..> end
- -> sub
/ -> div
+ -> add
* -> mul
[0.0000s]
th>
[0.0000s]
th> f = function() print('hello') end -- closures
[0.0000s]
th> f()
hello
[0.0000s]
th> a[f] = 'my function'
[0.0000s]
th> a
{
  function: 0x4059ed08 : "my function"
  + : "add"
  - : "sub"
  * : "mul"
  / : "div"
}
[0.0001s]
```

Torch tensors

Torch extends Lua's table with a Tensor object

- ▶ N-dimensional array
- ▶ Different types supported: IntTensor, FloatTensor, DoubleTensor
- ▶ torch package offers Matlab's common routines:

- ▶ zeros, ones, eye ...
- ▶ linear algebra stuff
- ▶ slice, view , narrow, fill

```
a = torch.Tensor(5,3) -- construct a 5x3 matrix, uninitialized
```

```
a = torch.rand(5,3)  
print(a)
```

```
b=torch.rand(3,4)
```

```
-- matrix-matrix multiplication: syntax 1  
a*b
```


```
-- matrix-matrix multiplication: syntax 2  
torch.mm(a,b)
```

```
-- matrix-matrix multiplication: syntax 3  
c=torch.Tensor(5,4)  
c:mm(a,b) -- store the result of a*b in c
```

- ▶ <https://github.com/torch/torch7/blob/master/doc/tensor.md>

Torch7 packages

- ▶ tensor handling
- ▶ neural networks
- ▶ optimization
- ▶ plotting
- ▶ statistics
- ▶ many more
- ▶ install via *luarocks*




Neural networks in Torch7 - nn package

Lego like base **Modules**:

- ▶ automatic differentiation
- ▶ each comes with 3 methods:
 - ▶ `output = forward(input)`
 - ▶ `gradInput = backward(input, gradOutput)`
computes the gradients of the module with respect to its own parameters, and its own inputs.
 - ▶ `zeroGradParameters()`
- ▶ internally keep two states variables: *output* and *gradInput*
- ▶ Linear, Convolution, Dropout, LookupTable, non linearities, etc

Neural networks in Torch7 - nn package


Complex networks combine modules using **Containers**:



```
th> model = nn.Sequential()
th> model:add( nn.Linear(10, 25) ) -- 10 input, 25 hidden units
th> model:add( nn.Tanh() ) -- some hyperbolic tangent transfer function
th> model:add( nn.Linear(25, 1) ) -- 1 output
th> print(model:forward(torch.randn(10)))
-0.3648
[torch.DoubleTensor of size 1]
```

Neural networks in Torch7 - nn package

Complex networks combine modules using **Containers**:



```
th> model = nn.Parallel(2,1)


th> model:add(nn.Linear(10,3))
nn.Parallel {
  input
  |'-> (1): nn.Linear(10 -> 3)
  ... -> output
}

th> model:add(nn.Linear(10,2))
nn.Parallel {
  input
  |'-> (1): nn.Linear(10 -> 3)
  |'-> (2): nn.Linear(10 -> 2)
  ... -> output
}

th> print(model:forward(torch.randn(10,2)))
-0.7266
 0.2804
 1.0107
 0.8189
-0.3478
[torch.DoubleTensor of size 5]
```

Neural networks in Torch7 - nn package

Complex networks combine modules using **Containers**:



```
th> model=nn.Concat(1);

th> model:add(nn.Linear(5,3))
nn.Concat {
  input
    |`-> (1): nn.Linear(5 -> 3)
    ... -> output
}

th> model:add(nn.Linear(5,7))
nn.Concat {
  input
    |`-> (1): nn.Linear(5 -> 3)
    |`-> (2): nn.Linear(5 -> 7)
    ... -> output
}

th> print(model:forward(torch.randn(5)))
-0.5925
 0.3837
-1.0290
-0.2023
-0.2447
 0.8625
-0.9372
-0.0472
-0.4417
 0.7487
[torch.DoubleTensor of size 10]
```

Neural networks in Torch7 - nn package

Complex networks combine modules using **Containers**:

- ▶ For table inputs, use `nn.ParallelTable()`, `nn.ConcatTable()`, `nn.SplitTable()`, `nn.JoinTable()`
- ▶ Arbitrary complex graphs can be made using these containers
 - ▶ Alternative: `nngraph` library

CIFAR-10 Example: Multi-layer Conv Net

- ▶ 3x32x32 input images
- ▶ 10 categories

airplane

automobile

bird

cat

deer


dog

frog

horse

ship

truck



```
net = nn.Sequential()
net.add(nn.SpatialConvolution(3, 6, 5, 5)) -- 3 input image channels, 6 output channels, 5x5 convolution kernel
net.add(nn.ReLU()) -- non-linearity
net.add(nn.SpatialMaxPooling(2,2,2,2)) -- A max-pooling operation that looks at 2x2 windows and finds the max.
net.add(nn.SpatialConvolution(6, 16, 5, 5))
net.add(nn.ReLU()) -- non-linearity
net.add(nn.SpatialMaxPooling(2,2,2,2))
net.add(nn.View(16*5*5)) -- reshapes from a 3D tensor of 16x5x5 into 1D tensor
net.add(nn.Linear(16*5*5, 120)) -- fully connected layer (matrix multiplication between input and weights)
net.add(nn.ReLU()) -- non-linearity
net.add(nn.Linear(120, 84)) -- non-linearity
net.add(nn.Linear(84, 10)) -- 10 is the number of outputs of the network (in this case, 10 digits)
net.add(nn.LogSoftMax()) -- converts the output to a log-probability. Useful for classification problems
```

```
input = torch.rand(1,32,32) -- pass a random tensor as input to the network
```


```
output = net.forward(input)
```

```
print(output)
```

```
net.zeroGradParameters() -- zero the internal gradient buffers of the network (will come to this later)
```

```
gradInput = net.backward(input, torch.rand(10))
```

```
print(#gradInput)
```



Loss functions - Criteria

- ▶ implemented just like neural network modules
- ▶ automatic differentiation
- ▶ two functions:
 - ▶ `forward(input, target)`
 - ▶ `backward(input, target)`
- ▶ negative log likelihood, max-margin, binary cross entropy, ...
- ▶ On CIFAR example:

```
criterion = nn.ClassNLLCriterion() -- a negative log-likelihood criterion for multi-class classification
criterion.forward(output, 3) -- let's say the groundtruth was class number: 3
gradients = criterion.backward(output, 3)
```

```
gradInput = net.backward(input, gradients)
```

Training

→ step 4/5: define a closure that estimates $f(x)$ and df/dx stochastically

```
08 -- define a closure, that computes the loss, and dloss/dx
09 feval = function()
10     -- select a new training sample
11     _nidx_ = (_nidx_ or 0) + 1
12     if _nidx_ > (#data)[1] then _nidx_ = 1 end
13
14     local sample = data[_nidx_]
15     local inputs = sample[1]
16     local target = sample[2]
17
18     -- reset gradients (gradients are always accumulated,
19     --                      to accomodate batch methods)
20     dl_dx:zero()
21
22     -- evaluate the loss function and its derivative wrt x,
23     -- for that sample
24     local loss_x = criterion:forward(model:forward(inputs), target)
25     model:backward(inputs, criterion:backward(model.output, target))
26
27     -- return loss(x) and dloss/dx
28     return loss_x, dl_dx
29 end
30
```

Source: <https://github.com/soumith/cvpr2015/>

Training / Optimization

- ▶ Optim package - many methods, easy plug-and-play

➔ **step 5/5: estimate parameters (train the model), stochastically**

```
31 -- SGD parameters
32 sgd_params = {learningRate = 1e-3, learningRateDecay = 1e-4,
33               weightDecay = 0, momentum = 0}
34
35 -- train for a number of epochs
36 epochs = 1e2
37 for i = 1,epochs do
38     -- this variable is used to estimate the average loss
39     current_loss = 0
40
41     -- an epoch is a full loop over our training data
42     for i = 1,(#data)[1] do
43
44         -- one step of SGD optimization (steepest descent)
45         _,fs = optim.sgd(feval,x,sgd_params)
46
47         -- accumulate error
48         current_loss = current_loss + fs[1]
49     end
50
51     -- report average error on epoch
52     current_loss = current_loss / (#data)[1]
53     print(' current loss = ' .. current_loss)
54 end
```

Source: <https://github.com/soumith/cvpr2015/>

What about GPU ?

Easy! Only 4 operations:

- ▶ `require 'cunn'`
- ▶ `net = net:cuda()`
- ▶ `criterion = criterion:cuda()`
- ▶ `input = input:cuda()`

Everything else stays the same!

Resources and Tutorials

- ▶ Main website: `www.torch.ch`
- ▶ Torch cheat sheet:
`https://github.com/torch/torch7/wiki/Cheatsheet`
- ▶ Tutorials:
 - ▶ Fast th basics:
`http://rnduja.github.io/tags/torch/`
 - ▶ 60-minute blitz:
`https://github.com/soumith/cvpr2015`
 - ▶ MNIST digit classifier (how to properly use minibatches):
`https://github.com/torch/demos/blob/master/train-a-digit-classifier/train-on-mnist.lua`
 - ▶ Supervised, unsupervised, graphical models:
`https://github.com/torch/tutorials`