

Deep Learning: Theory and Practice

Tensorflow Practical Session

Overview

Material

<https://github.com/dalab/deep-learning-workshop-public.git>

Program

1. Session 1: Introduction & simple logistic regression
2. Session 2: Convolutional neural networks for image classification
3. Session 3: Recurrent and convolutional neural networks for text sentiment

Requirements for practical exercises

- Have tensorflow installed
- Have the datasets downloaded

Quick introduction

Now

1. Tensorflow basics
2. Constructing a simple model and training it
3. Monitoring, Printing, Debugging

Variables and Tensors in Tensorflow

Variables

- Maintain their state
- Different methods to create...

```
W1 = tf.Variable([[0.1,2.2], [-2.7,0.3]], name = "weights")
```

```
W2 = tf.Variable(tf.zeros([200, 200]), name = "weights")
```

```
W3 = tf.Variable(tf.random_uniform([200, 200], -10, 10), name = "weights")
```

Tensors

At every node of the graph, the result of a computation is a tensor.

Tensors have...

- ... a shape, e.g.
 - {} (scalar)
 - [20] (20-d vector)
 - [50,10] (50x10 matrix)
 - [32,3,100,100] (higher order tensors)
- ...a type e.g. `tf.int32`, `tf.float32`, `tf.bool`, `tf.string`

The graph concept in TF

- Variables

- (Input) Placeholders

```
x = tf.placeholder(tf.float32, [200,10], "input")
```

- Operations (e.g. for matrices A,B,C)

```
C = A - B
```

```
C = tf.matmul(A,B)
```

```
C = tf.nn.relu(A)
```

Common operations include

- Pointwise: `tf.mul`, `tf.add`, `tf.sigmoid`, `tf.tanh`,
`tf.nn.relu`
- Summing and averaging: `tf.reduce_sum`, `tf.reduce_mean`

Initialization and Sessions

Starting the session fixes the graph and places the ops on devices.

```
x = tf.placeholder(tf.float32, [200,10], "input") # Input placeholder
W_1 = tf.Variable(tf.zeros([200, 200]), name = "weights")
hidden_1 = ... # Some deep network here
hidden_2 = ...

...
y = tf....      # Network output

# Run graph to get output given an input
with tf.Session() as session:
    datapoint = ... # some np array of size [200, 10]
    feed_dict = {x : datapoint} # Map TF placeholders to numpy arrays
    y_output = session.run(y, feed_dict = feed_dict) #run the graph

    print("Output is " + y_output)
```

Training

Example:

```
loss = ...    # some scalar tensor
opt = tf.train.GradientDescentOptimizer(0.01) # here fixed learning rate
update_step = opt.minimize(loss, global_step)
```

See also


https://www.tensorflow.org/versions/r0.9/api_docs/python/train.html

Example: Autoencoder

The Problem

Given data $\mathcal{X} = \{x_1, \dots, x_n\} \subset \mathbb{R}^d$
 and $\tilde{d} < d$, find $f : \mathbb{R}^d \rightarrow \mathbb{R}^{\tilde{d}}$ and $g : \mathbb{R}^{\tilde{d}} \rightarrow \mathbb{R}^d$
 so that for $h_i = f(x_i) \in \mathbb{R}^{\tilde{d}}$
 the reconstruction error $\sum_i \|x_i - g(h_i)\|_2^2$ is small

Neural network approach with single hidden layer



$$h = f(x) = \sigma(\mathbf{W}_{\text{enc}}x + b_{\text{enc}}) \quad \text{with } \mathbf{W}_{\text{enc}} \in \mathbb{R}^{\tilde{d} \times d}, b_{\text{enc}} \in \mathbb{R}^{\tilde{d}}$$

$$y = g(h) = \sigma(\mathbf{W}_{\text{dec}}h + b_{\text{dec}}) \quad \text{with } \mathbf{W}_{\text{dec}} \in \mathbb{R}^{d \times \tilde{d}}, b_{\text{dec}} \in \mathbb{R}^d$$

Example: Autoencoder in TF

```
d_data = 100
d_hidden = 30

# Construct Graph
x = tf.placeholder(tf.float32, [d_data,1], name="input")


# Hidden Layer Variables
W_enc = tf.Variable(tf.random_uniform([d_hidden,d_data], -1, 1), name="W_enc")
b_enc = tf.Variable(tf.zeros([d_hidden,1]), name = "bias_enc")

W_dec = tf.Variable(tf.random_uniform([d_data,d_hidden], -1, 1), name="W_dec")
b_dec = tf.Variable(tf.zeros([d_data]), name = "bias_dec")

# Hidden layer graph
h = (tf.matmul(W_enc, x) + b_enc)

# Output and reconstruction loss
y = (tf.matmul(W_dec, h) + b_dec)
loss = tf.sqrt(tf.reduce_sum(tf.square(x - y)))

# Optimizer
opt = tf.train.GradientDescentOptimizer(0.01)
update_step = opt.minimize(loss)
```



Printing and Debugging

Printing

You cannot access a tensor's content e.g. `W[0][1]`

Only properties such as `W.get_shape()` or `W.type` are available.

- Option 1, one-time printing: Run the tensor in a session

```
W_data = session.run(W)    #get numpy.ndarray
print(W_data)
```

- Option 2, print continuously: Print node

```
W = tf.Variable(tf.zeros([100,100]),...)
W = tf.Print(W, [W], message="Entries of W: ")
X = tf.matmul(W, X)... # W is still a matrix
```

Debugging the graph

- In tensorboard manually check all dependencies in the graph

See also

https://www.tensorflow.org/versions/r0.9/get_started/basic_usage.html

Example: Autoencoder in TF

```
data = np.random.rand(d_data, n_datapoints + 1)

# Train
with tf.Session() as session:
    # Initialize variables
    init = tf.initialize_all_variables()
    session.run(init)


    # Do nsteps many SGD update steps
    for i in range(10000):
        datapoint = data[:, np.random.randint(0, n_datapoints)]
        feed_dict = {x : np.transpose([datapoint])}
        session.run(update_step, feed_dict = feed_dict)

    # Every now and then test the loss on our hold out datapoint
    if i % 200 == 0:
        test_point = data[:, n_datapoints]
        feed_dict = {x : np.transpose([test_point])}
        test_loss = session.run(loss, feed_dict = feed_dict)
        print("test loss is %f " % test_loss)
```

Tensorboard

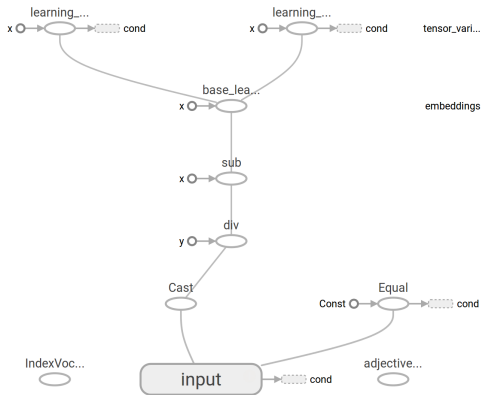

test_loss

test_loss



train_loss

train_loss



Tensorboard

How to add a summary

For example, track the norm of $W = \text{tf.Variable}(\text{tf.zeros}([100,100]), \dots)$

At graph construction time

Once globally:

- Create a summary writer
`summary_writer = tf.train.SummaryWriter("/my/directory")`

For this particular summary:

- Get the norm: `W_norm = tf.sqrt(tf.sum_reduce(tf.square(W)))`
- Create summary: `W_summary = tf.scalar_summary("Norm of W", W_norm)`

At training time

- Merge all summaries: `summaries = tf.merge_summary([W_summary, ...])`
- Fetch the data and turn into a formatted string:
`summary_str = session.run(summaries)`
- Send the string to the writer
`summary_writer.add_summary(summary_str, global_step)`

Resources

The *How to* and *Tutorials* section on tensorflow.org are actually good resources to recap concepts.

https://www.tensorflow.org/versions/r0.9/get_started/index.html

<https://www.tensorflow.org/versions/r0.9/tutorials/index.html>

For specific problems, google search often delivers quite useful threads on

<http://stackoverflow.com> and

<https://github.com/tensorflow/tensorflow/issues>