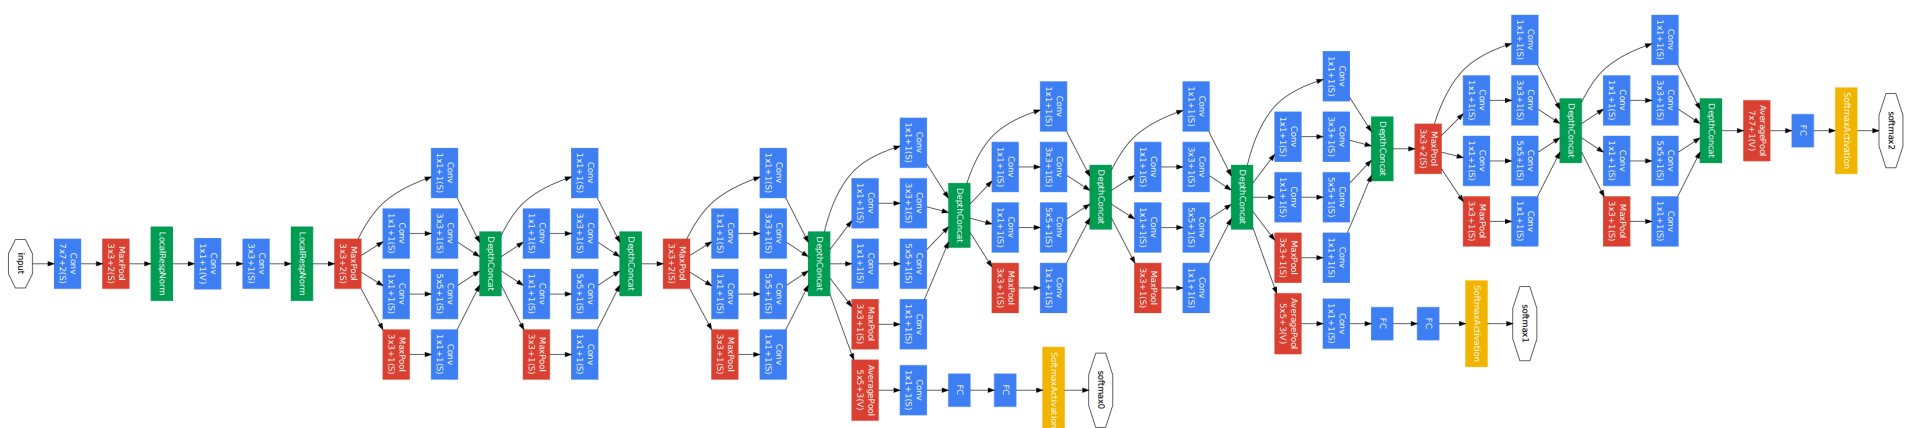


Theano & Tensorflow

Symbolic Compute Graphs


GoogLeNet





Compute Graphs


- Deep networks are **static**
- Math can be represented as a **graph**

z := x + y






$$z := a * x + b * y + c$$



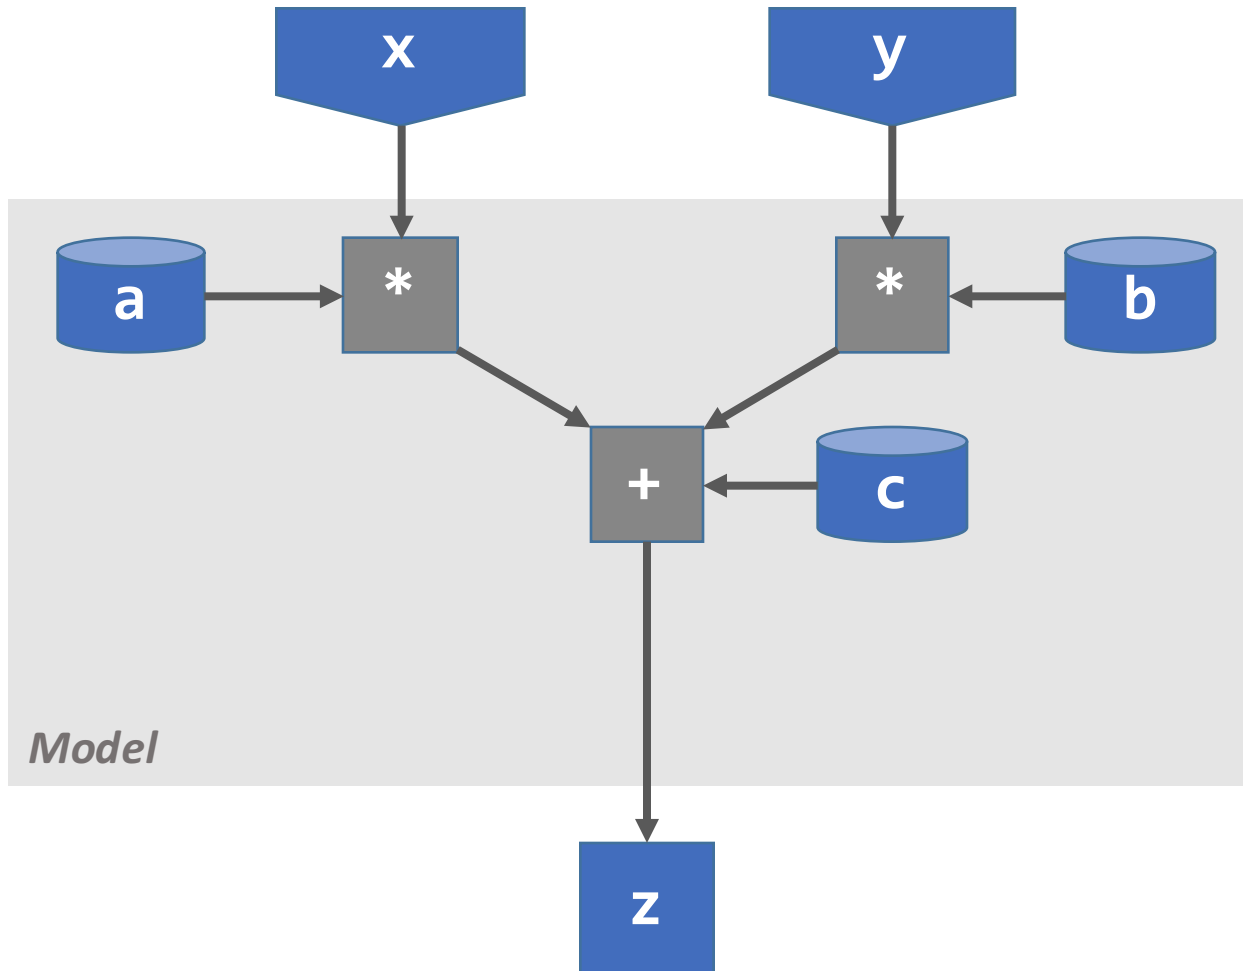
$z := a * x + b * y + c$

Parameters




$$z := a * x + b * y + c$$

Parameters

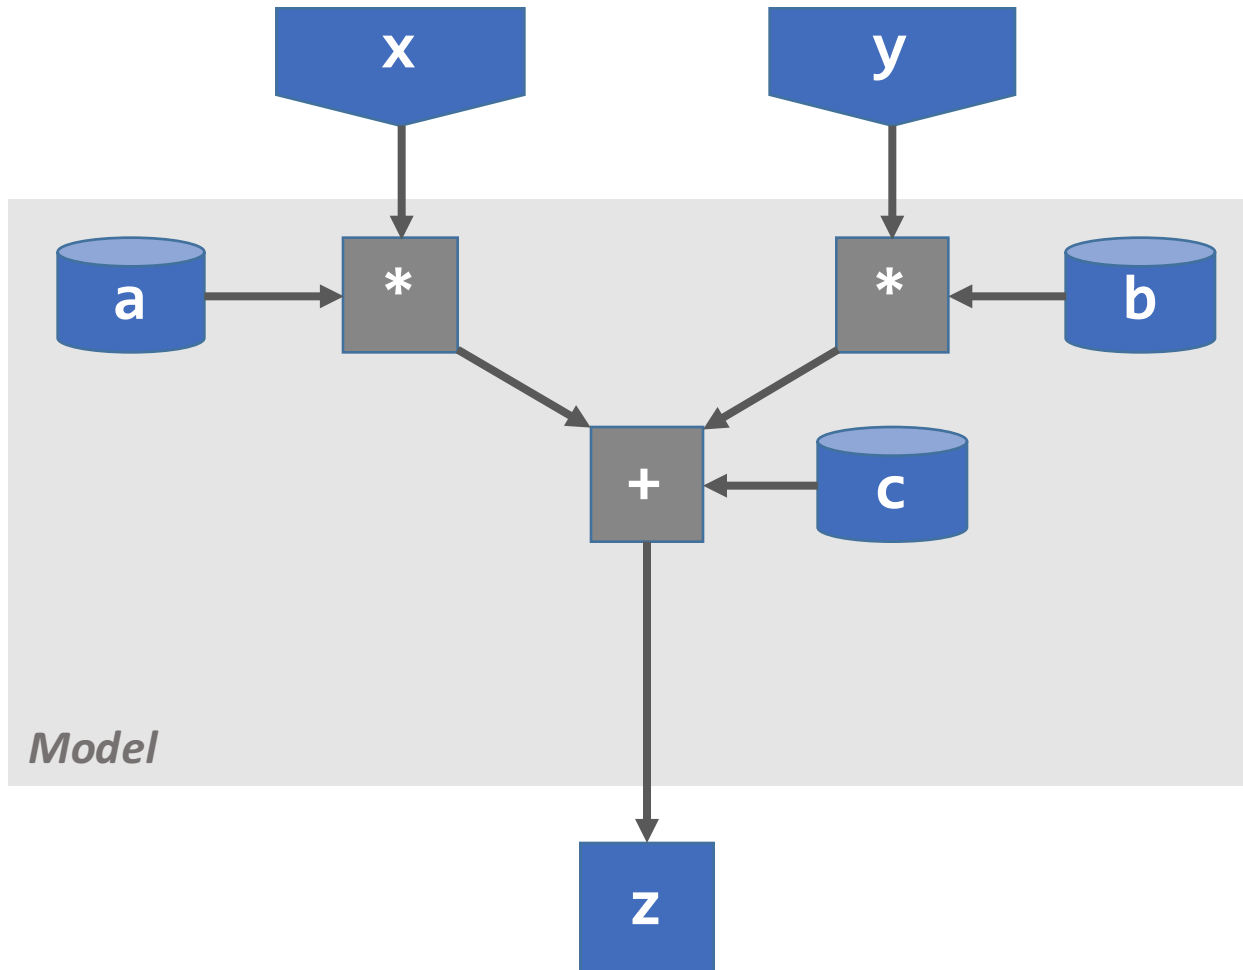


$z := \text{Exp}(a * x + b * y + c)$



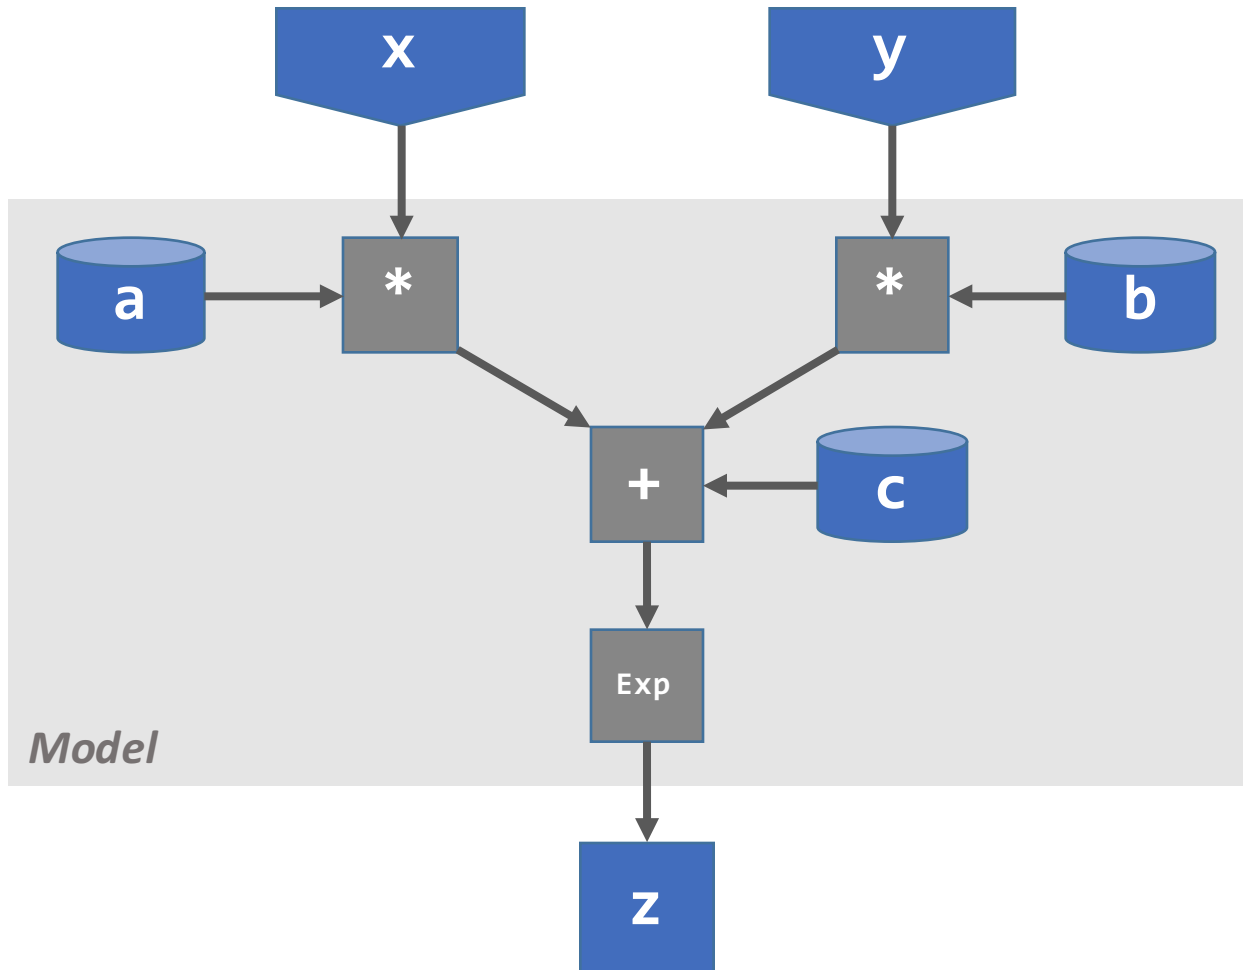
$z := \text{Exp}(a * x + b * y + c)$

Function

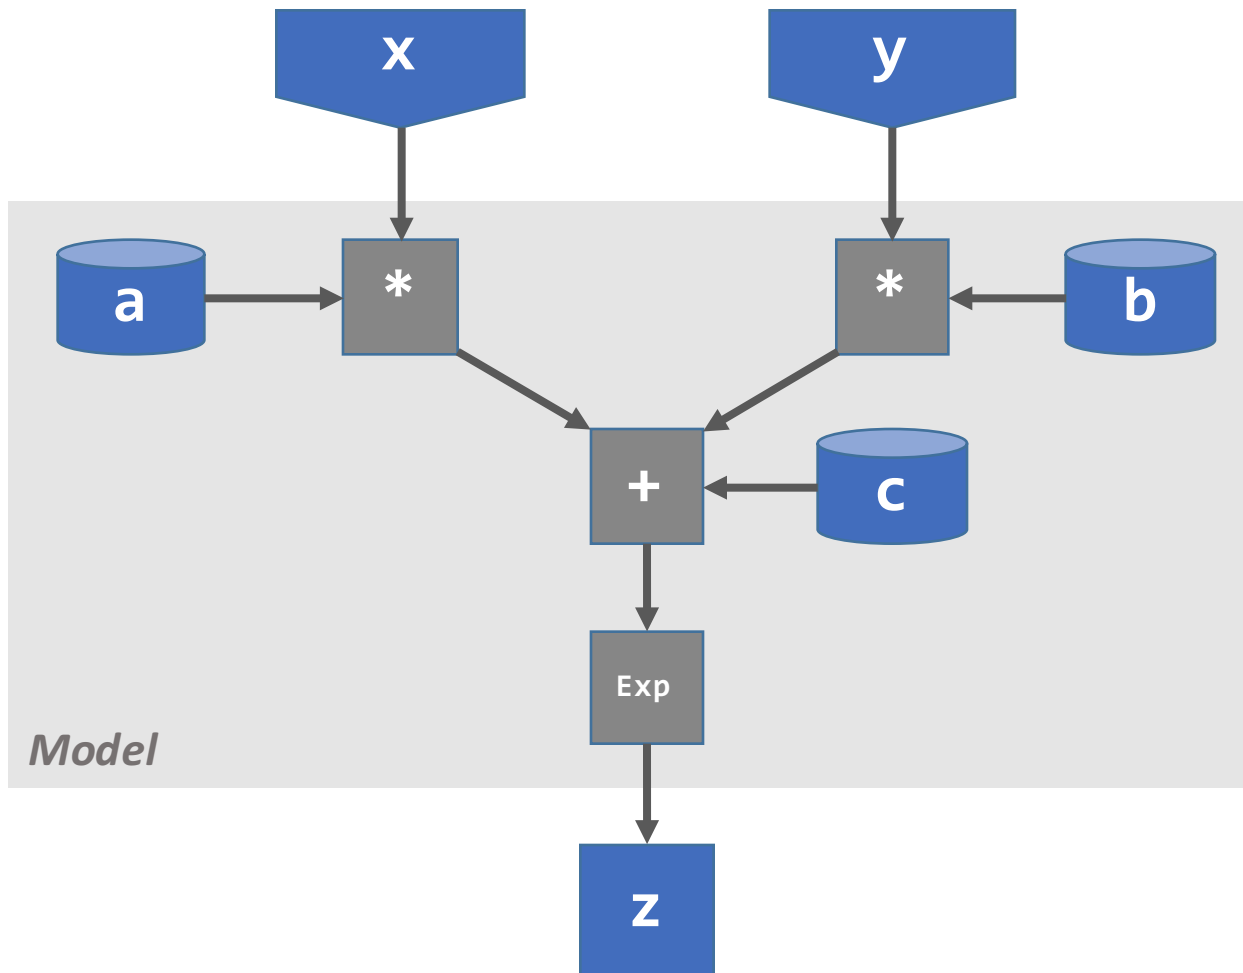


$z := \text{Exp}(a * x + b * y + c)$

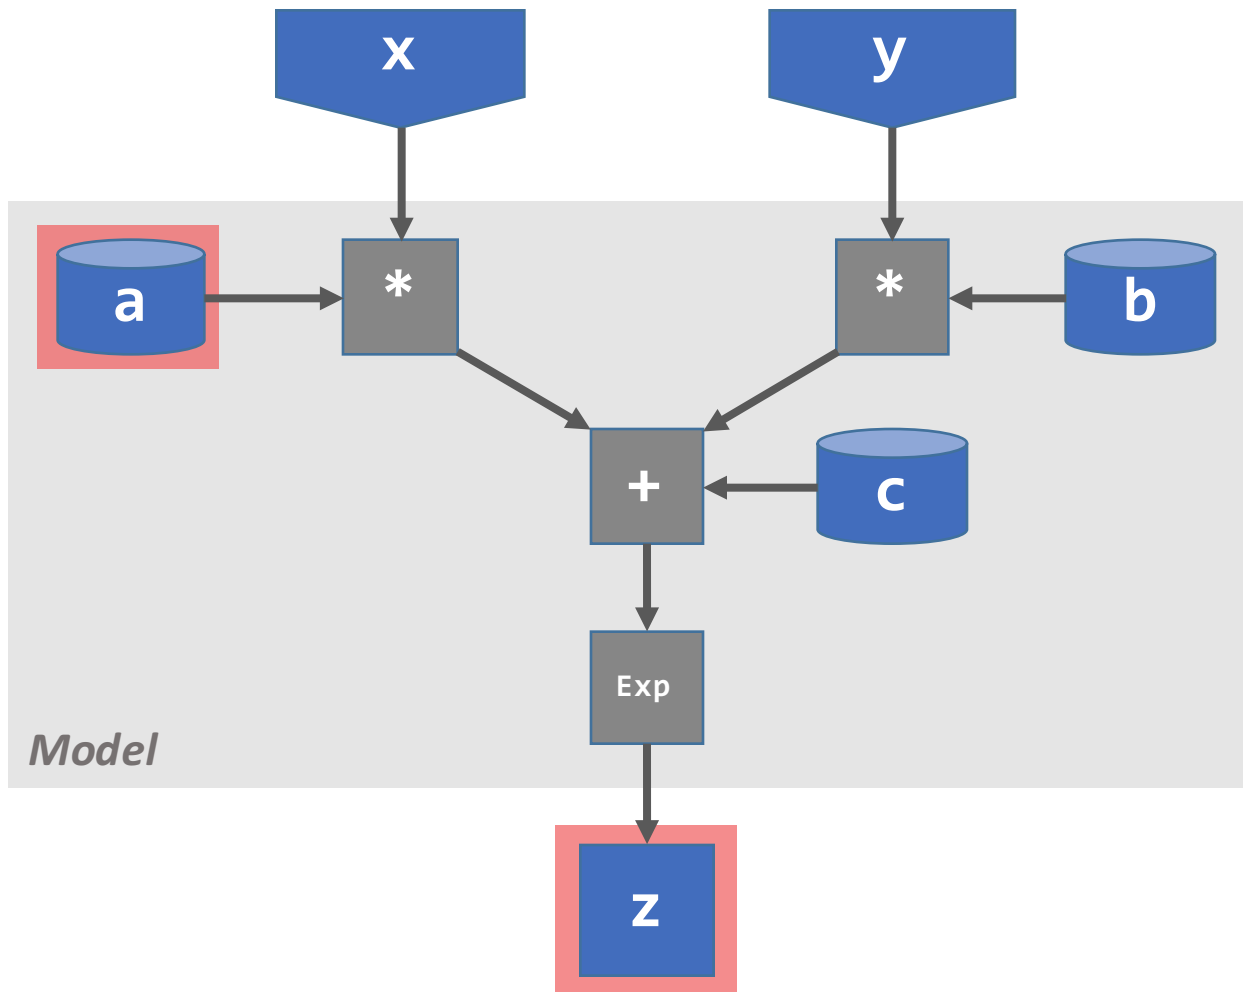
Function



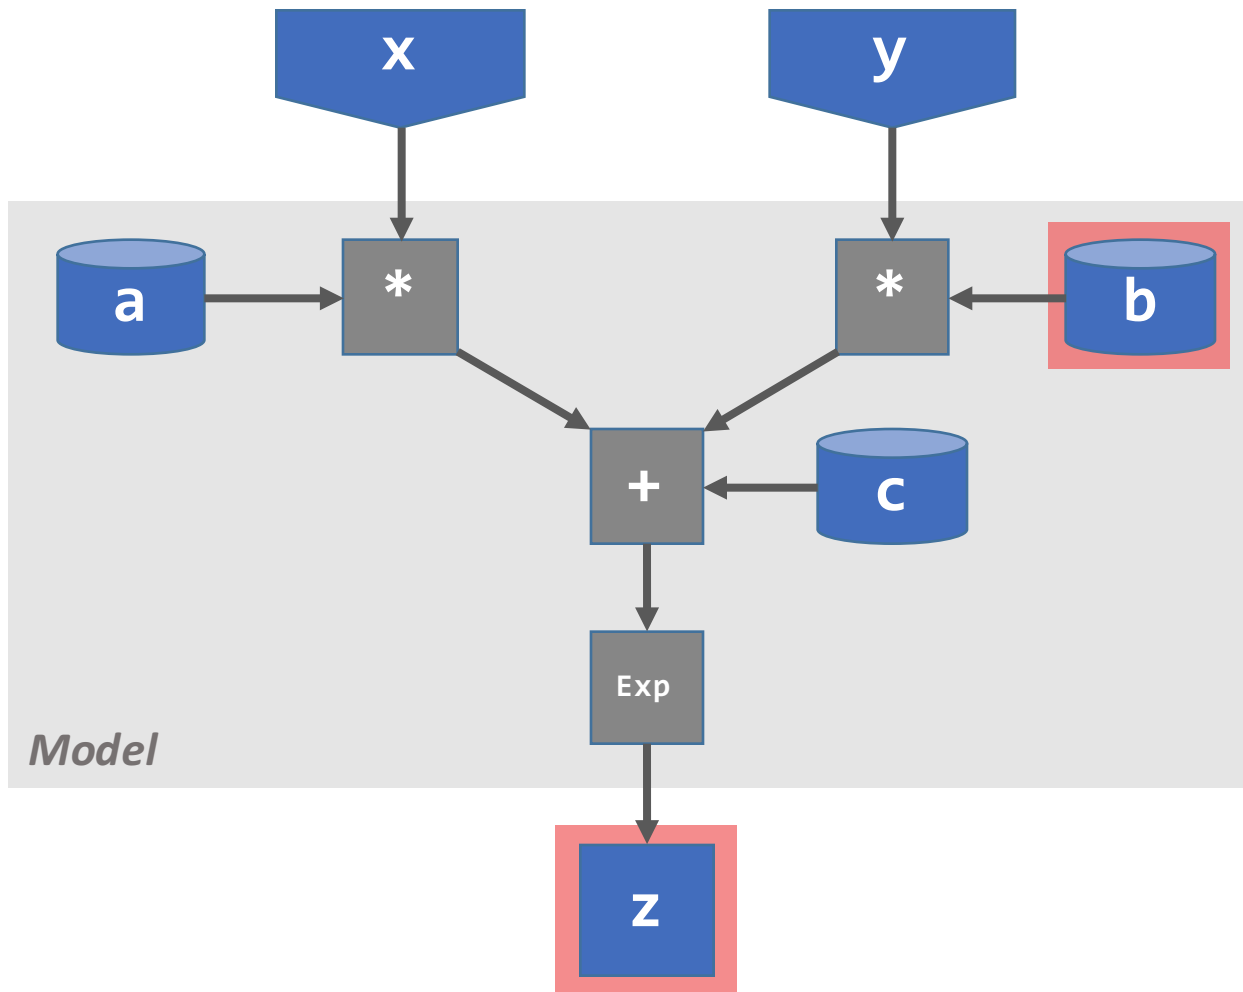
What is $\partial z / \partial a$?



What is $\partial z / \partial a$?

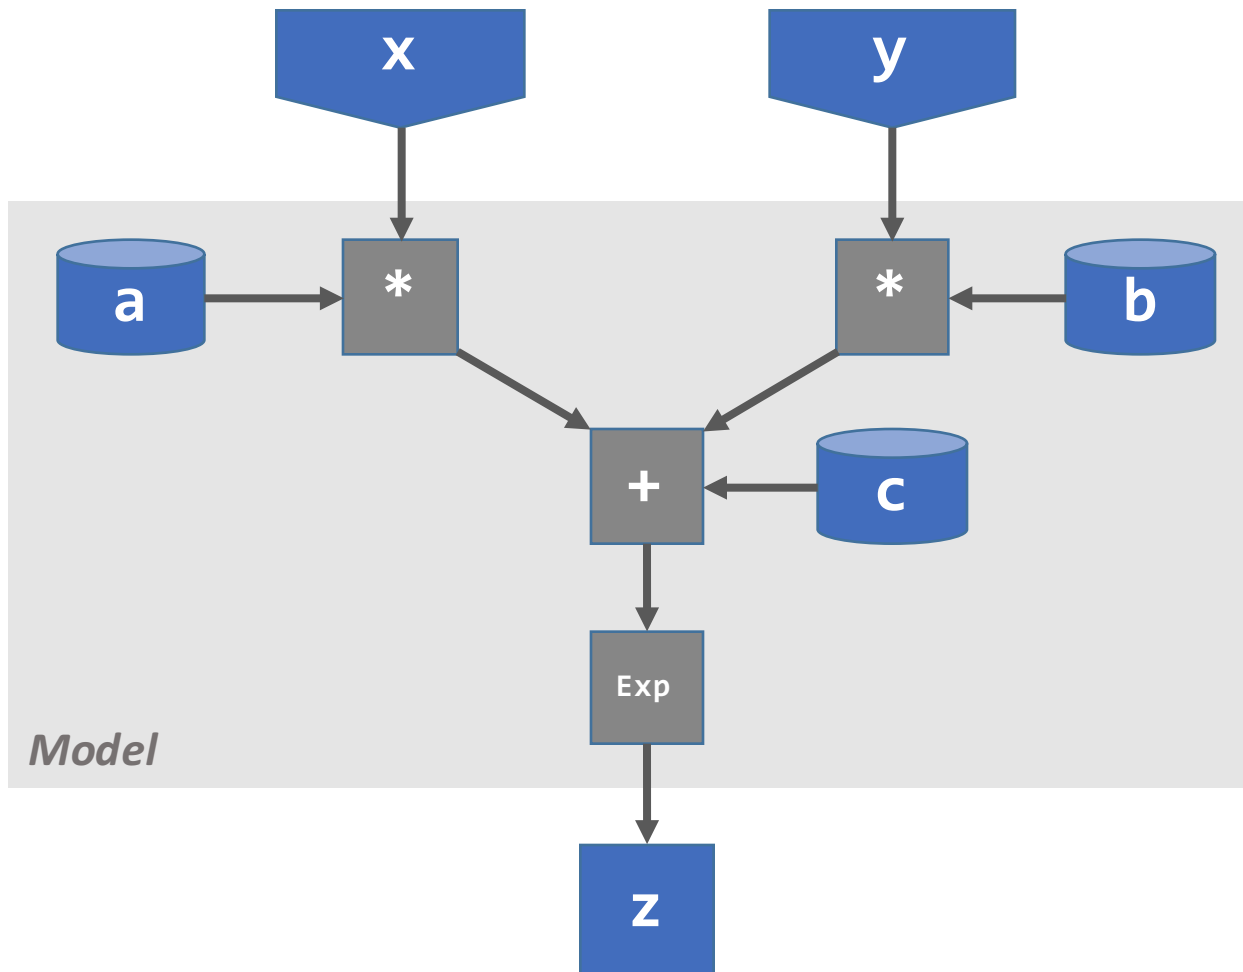


What is $\partial z / \partial a$? $\partial z / \partial y$?

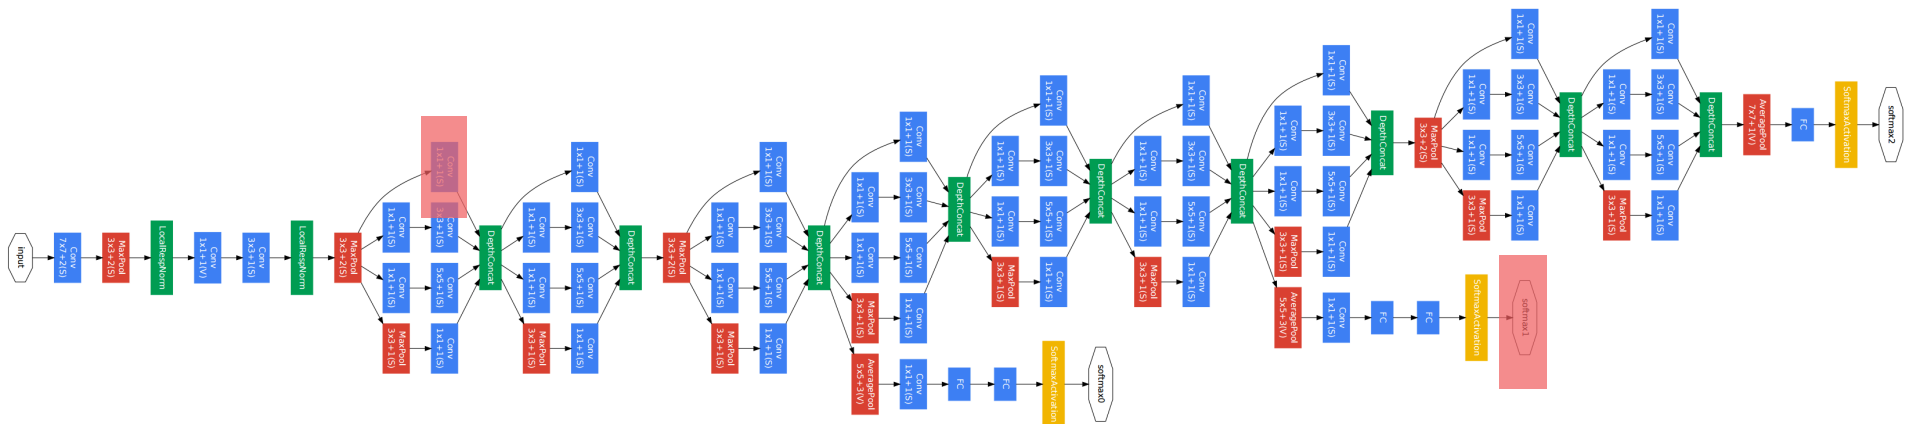


What is $\partial z / \partial a$? $\partial z / \partial y$?

How about $\partial \oplus / \partial b$?

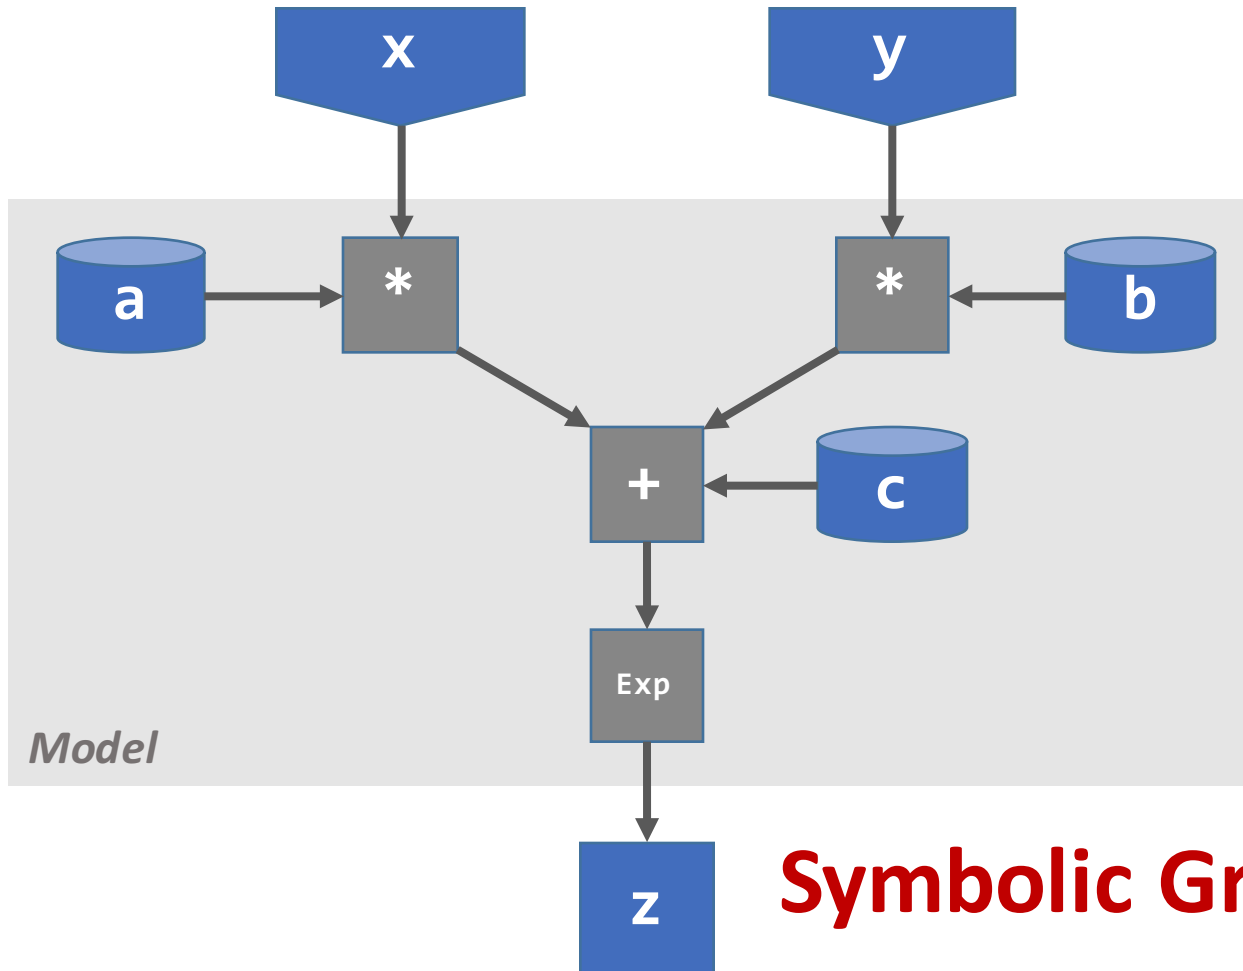


How about here?

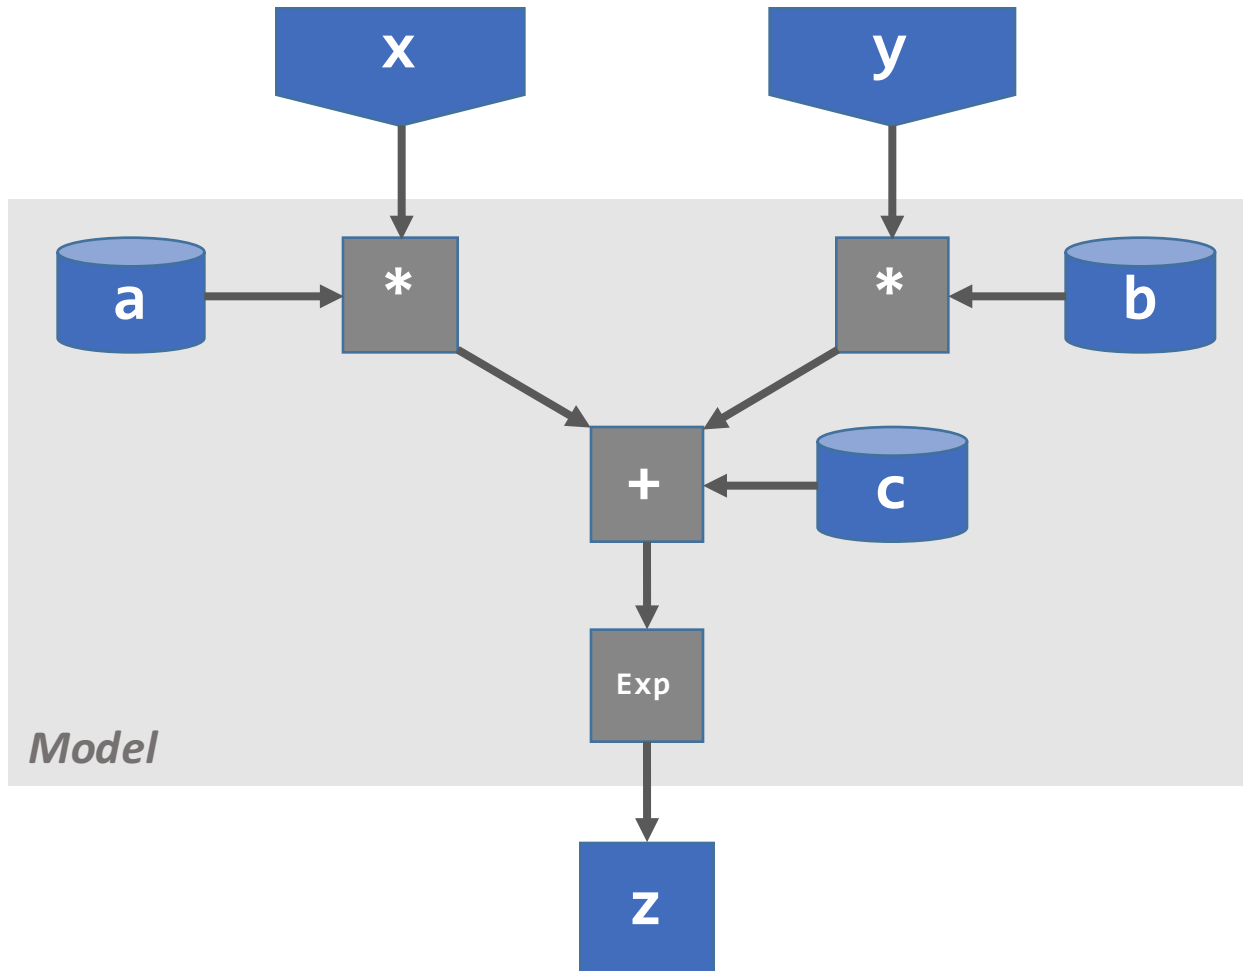


What is $\partial z / \partial a$? $\partial z / \partial y$?


How about $\partial \oplus / \partial b$?



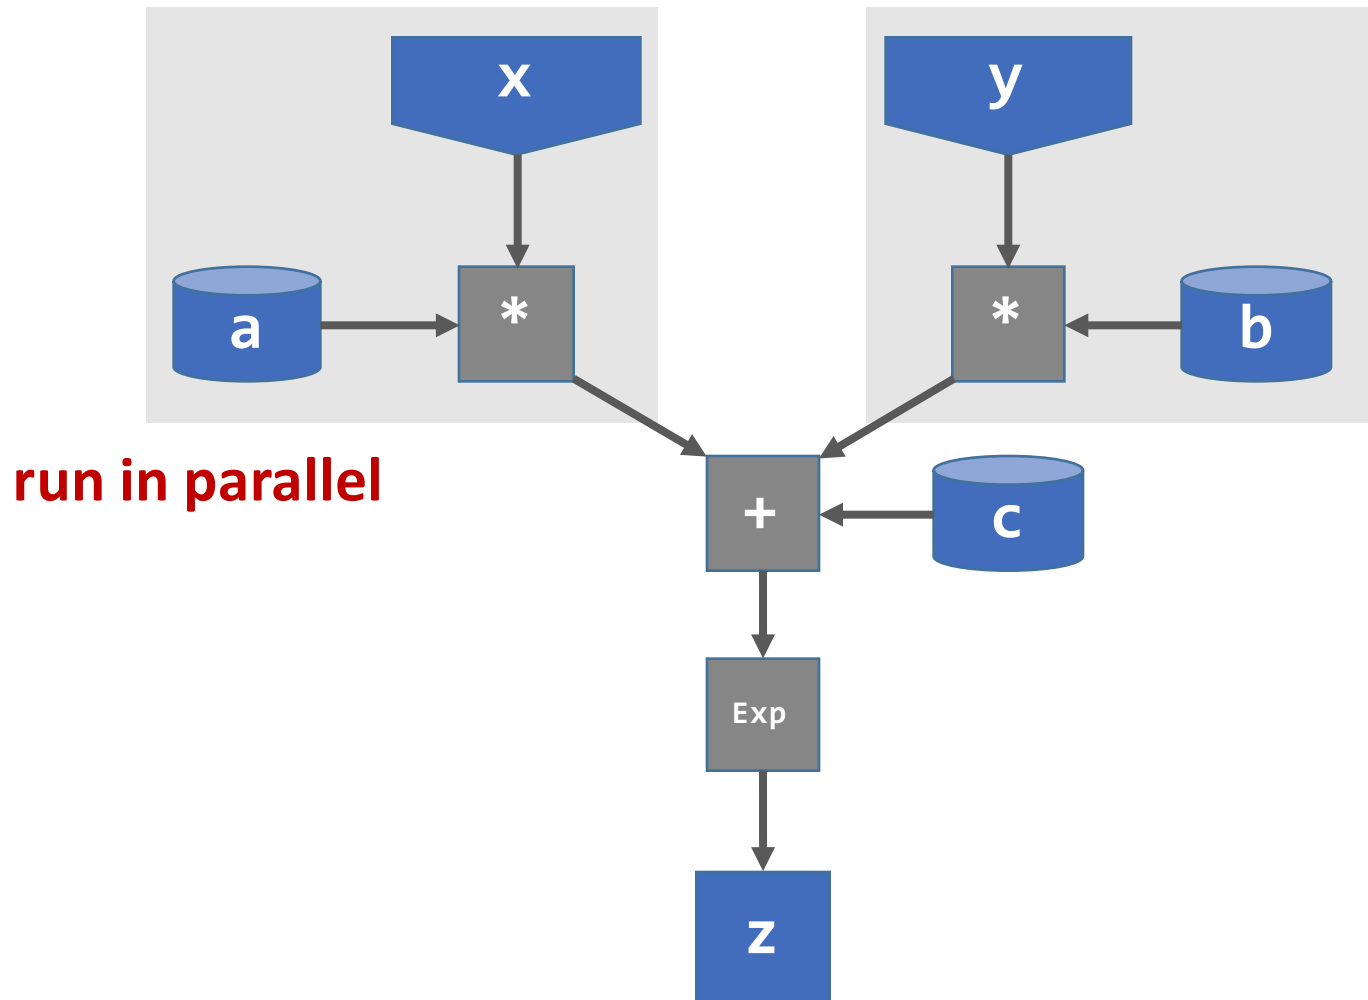
What about optimization?




What about optimization?



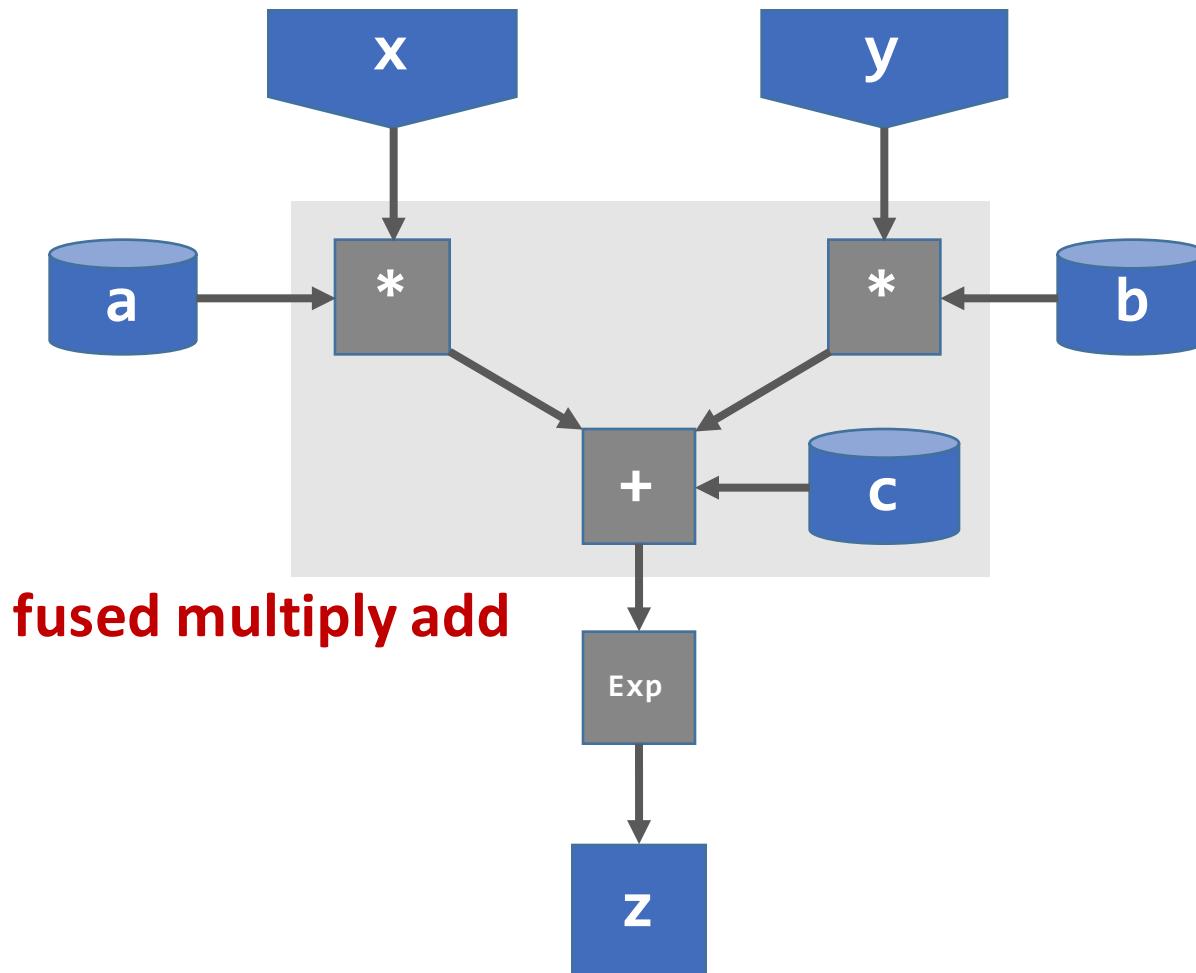
What about optimization?



What about optimization?



What about optimization?



Compute Graphs

- Deep networks are **static**
- Math can be represented as a **graph**

Compute Graphs

- Deep networks are **static**
- Math can be represented as a **graph**
- This allows to automate
 - Differentiation
 - Optimization
 - Parallelization

How do I define a graph?

How do I define a graph?

- You don't have to worry about this!
- Mostly transparent to the user
- Math abstracted into **network layers**

```
X, Y = matrix(), vector()
```

Inputs

```
X, Y = matrix(), vector()
```

```
net = InputLayer(X)  
net = DenseLayer(net, num_units=64)  
net = DenseLayer(net, num_units=10,  
                  nonlinearity=softmax)
```

```
prediction = get_output(net)  
parameters = get_all_params(net)
```

Network

```
X, Y = matrix(), vector()
```

```
net = InputLayer(X)
```

```
net = DenseLayer(net, num_units=64)
```

```
net = DenseLayer(net, num_units=10,  
                  nonlinearity=softmax)
```

```
prediction = get_output(net)
```

```
parameters = get_all_params(net)
```

```
loss = mean(categorical_crossentropy(prediction, Y))
```

```
accuracy = mean(eq(prediction, Y))
```

Cost

```
X, Y = matrix(), vector()
```

```
net = InputLayer(X)
```

```
net = DenseLayer(net, num_units=64)
```

```
net = DenseLayer(net, num_units=10,  
                  nonlinearity=softmax)
```

```
prediction = get_output(net)
```

```
parameters = get_all_params(net)
```

```
loss = mean(categorical_crossentropy(prediction, Y))
```

```
accuracy = mean(eq(prediction, Y))
```

```
gradient = grad(loss, parameters)
```

```
updates = sgd(grad, parameters)
```

Optimization

```
X, Y = matrix(), vector()
```

```
net = InputLayer(X)
```

```
net = DenseLayer(net, num_units=64)
```

```
net = DenseLayer(net, num_units=10,  
                  nonlinearity=softmax)
```

```
prediction = get_output(net)
```

```
parameters = get_all_params(net)
```

```
loss = mean(categorical_crossentropy(prediction, Y))
```

```
accuracy = mean(eq(prediction, Y))
```

```
gradient = grad(loss, parameters)
```

```
updates = sgd(grad, parameters)
```

```
f_train = theano.function([X, Y], loss, updates=updates)
```

```
f_valid = theano.function([X, Y], [loss, accuracy])
```

```
f_predict = theano.function([X], prediction)
```

Compile


```
f_train = theano.function([X, Y], loss, updates=updates)
f_valid = theano.function([X, Y], [loss, accuracy])
f_predict = theano.function([X], prediction)
```

```
{{ train loop }}
```

```
x_mb, y_mb = get_minibatch()
```

```
f_train(x_mb, y_mb)
```

```
{{ end }}
```



Numpy Arrays

Train

```
f_train = theano.function([X, Y], loss, updates=updates)
f_valid = theano.function([X, Y], [loss, accuracy])
f_predict = theano.function([X], prediction)
```

```
{{ train loop }}
```

```
x_mb, y_mb = get_minibatch()
f_train(x_mb, y_mb)
```

```
{{ end }}
```

```
x_v, y_v = get_validation_set()
valid_loss, valid_acc = f_valid(x_v, y_v)
```

Validation

```
f_train = theano.function([X, Y], loss, updates=updates)
f_valid = theano.function([X, Y], [loss, accuracy])
f_predict = theano.function([X], prediction)
```

```
{{ train loop }}
```

```
x_mb, y_mb = get_minibatch()
f_train(x_mb, y_mb)
```

```
{{ end }}
```

```
x_v, y_v = get_validation_set()
valid_loss, valid_acc = f_valid(x_v, y_v)
```

```
x_new = get_unlabeled_data()
y_new = f_pred(x_new)
```

Prediction

```
f_train = theano.function([X, Y], loss, updates=updates)
f_valid = theano.function([X, Y], [loss, accuracy])
f_predict = theano.function([X], prediction)
```

```
{{ train loop }}
```

```
x_mb, y_mb = get_minibatch()
f_train(x_mb, y_mb)
```

```
{{ end }}
```

```
x_v, y_v = get_validation_set()
valid_loss, valid_acc = f_valid(x_v, y_v)
```

```
x_new = get_unlabeled_data()
y_new = f_pred(x_new)
```

Prediction

It's just python!

Pros & Cons

- + Automatic differentiation & optimization**
- + Transparent GPU usage**
- + Programming in Python**
- Only Static Graphs**
- Learning Curve**

Theano or Tensorflow?

- **Theano**

- Made by researchers, for researchers
- Mature
- Compiles C on the fly

- **Tensorflow**

- Made by Google
- Still in Beta
- Focus on large scale learning & production

Links

- **Theano + Lasagne**

- `github.com/theano/theano`
- `github.com/lasagne/lasagne`

- **Tensorflow**

- `github.com/tensorflow/tensorflow`

- **Keras**

- `github.com/fchollet/keras`